


How an LLM-generated CARE skill becomes an auditable acquisition rule

The LLM proposes a bounded patch once. Target observations calibrate evidence; Python executes the equations and gate deterministically.



If evidence is weak or acquisition loss exceeds the round-dependent risk budget, the selected candidate is the target-only anchor.