



Can the LLM estimate whether its next experiment will improve the current best?

170 archived decisions / 17 routes; observed improvement rate 0.159. Brier: LLM 0.133, same-candidate GP 0.134. Equal-route difference -0.002 [-0.013, +0.009].



Real archived decisions; no generated values. Probabilities are scored on the executed candidate. Retrospective calibration evidence, not prospective gate validation.