


# Prospective FLIP2 IRED confirmation: an outcome-blind LLM skill transfers positively

The LLM hypothesis, executable additive skill, gate rule, code hashes, and seed schedules were frozen before the official file was downloaded.

- RGPE
- Fixed skill prior
- ICM-BMA
- LLM additive skill

## A. Positive-transfer effect and strong baselines



Interpretation: the preregistered LLM additive skill improves AUC versus target-only GP, but the fixed skill prior is stronger; this is one-family finite-pool evidence, not universal transfer or wet-lab validation.

## B. Untouched official-test trajectories

